




JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING


Introduction to Machine Learning

k-Means and *k*-medoids Clustering

Some Data



Grid Clustering



Expectation Maximization

A local search process where we estimate parameters and then adjust them to increase likelihood of correctness

Example – Find one of the roots of $x^5 - 3x^2 + 2x - 17 = 0$.

- Rewrite as $x = (3x^2 - 2x + 17)^{1/5}$.
- "Guess" that $x = 1$.
- Substitute into right hand side and recalculate, $x = 1.7826$.
- Repeat until convergence

$1 \rightarrow 1.7826 \rightarrow 1.8716 \rightarrow 1.8845 \rightarrow 1.8864 \rightarrow 1.8866 \rightarrow 1.8867 \rightarrow \dots$

The proof sets up what we call a "contraction map."

EM for Clustering

- Bivariate Gaussians with equal variance
- k -means clustering
- **Assumptions**
 - We have k multi-variate Gaussian/spherical clusters.
 - Each cluster has an unknown mean vector: $\langle \mu_1, \dots, \mu_k \rangle$.
 - We do not know which Gaussian generated which data point.
- We can apply EM to find the cluster means.
- We can use the cluster means to assign the points.

***k*-Means Clustering**

Example

- Let $k = 2$.
- Describe a generated instance as $y_i = \langle x_i, z_{i1}, z_{i2} \rangle$, where $z_{ij} = 1$ means x_i was generated by cluster j .
- Begin by guessing values for $h = \langle \mu_1, \mu_2 \rangle$ at random.

E Step: (Expectation)

$$\begin{aligned} E[z_{ij}] &= \frac{P(x_i | \mu_j)}{\sum_{k=1}^2 P(x_i | \mu_k)} \\ &= \frac{\exp\left[-\frac{1}{2\sigma^2}(x_i - \mu_j)^2\right]}{\sum_{k=1}^2 \exp\left[-\frac{1}{2\sigma^2}(x_i - \mu_k)^2\right]} \end{aligned}$$

M Step: (Maximization, $h' = \langle \mu'_1, \mu'_2 \rangle$)


$$\mu_j = \frac{\sum_{i=1}^m E[z_{ij}] x_i}{\sum_{i=1}^m E[z_{ij}]}$$

k-Means Clustering

Algorithm 10.3 *K*-Means Clustering

```
1: function KMEANS( $\mathcal{D}, k$ )
2:   initialize  $\mu_1, \dots, \mu_k$  randomly
3:   repeat
4:     for all  $\mathbf{x}_i \in \mathcal{D}$  do
5:        $c \leftarrow \arg \min_{\mu_j} d(\mathbf{x}_i, \mu_j)$             $\triangleright d()$  is the distance between  $\mathbf{x}_i$  and  $\mu_j$ .
6:       assign  $\mathbf{x}_i$  to the cluster  $c$ 
7:     end for
8:     recalculate all  $\mu_j$  based on new clusters
9:   until no change in  $\mu_1, \dots, \mu_k$ 
10:  return  $\mu_1, \dots, \mu_k$ 
11: end function
```

Illustrating *k*-Means



Another Illustration of k -Means

1 2 3 5 6 7 10

$K = 2$

1 2 3 5
└──┬──┘

2.5

6 7 10
└──┬──┘

7.67

1 2 3
└──┬──┘

2

3 6 7 10
└──┬──┘

7.5

K-Medoids

- Representative object or member of a data set
- Different objective functions
 - *K*-means

$$J(\mathcal{C}) = \sum_{j=1}^K \sum_{x_i \in \mathcal{C}} (x_i - c_j)^2$$

- *K*-medoids

$$J(\mathcal{C}) = \sum_{j=1}^K \sum_{x_i \in \mathcal{C}} (x_i - m_j)^2$$

Partitioning Around Medoids (PAM)

Algorithm 10.4 K -Medoids Clustering

```
1: function KMEDOIDS-PAM( $\mathcal{D}, k$ )
2:   select  $\mathbf{m}_1, \dots, \mathbf{m}_k$  randomly
3:   repeat
4:     for all  $\mathbf{x}_i \in \mathcal{D}$  do
5:        $c \leftarrow \arg \min_{\mathbf{m}_j} d(\mathbf{x}_i, \mathbf{m}_j)$   $\triangleright d()$  is the distance between  $\mathbf{x}_i$  and  $\mathbf{m}_j$ .
6:       assign  $\mathbf{x}_i$  to the cluster  $c$ 
7:     end for
8:      $distortion_{k\text{-medoids}} = \sum_{j=1}^k \sum_{i \in \text{ownedby}(\mathbf{c}_j)} (\mathbf{x}_i - \mathbf{m}_j)^2$ 
9:     for all  $\mathbf{m}_i \in \mathbf{m}$  do
10:      for all  $\mathbf{x}_j \in \mathcal{D}$  where  $\mathbf{x}_j \notin \mathbf{m}$  do
11:        swap  $\mathbf{m}_i$  and  $\mathbf{x}_j$ 
12:         $distortion'_{k\text{-medoids}} = \sum_{j=1}^k \sum_{i \in \text{ownedby}(\mathbf{c}_j)} (\mathbf{x}_i - \mathbf{m}_j)^2$ 
13:        if  $distortion_{k\text{-medoids}} \leq distortion'_{k\text{-medoids}}$  then
14:          swap back
15:        end if
16:      end for
17:    end for
18:  until no change in  $\mathbf{m}_1, \dots, \mathbf{m}_k$ 
19:  return  $\mathbf{m}_1, \dots, \mathbf{m}_k$ 
20: end function
```

Fuzzy c -Means

- A method to create “soft” clusters, where f is a level of “fuzzification” in the range $1 \dots n$
 - $f = 1$ indicates “crisp” clusters
 - $f > 1$ indicates “soft” clusters
- Objective function

$$J(\mathcal{C}) = \sum_{i=1}^n \sum_{j=1}^c w_j(x_i)^f \|c_j - x_i\|^2$$

$$w_j(x_i) = \frac{1}{\sum_{k=1}^c \left(\frac{\|c_j - x_i\|}{\|c_k - x_i\|} \right)^{2/(f-1)}}$$

$$c_j = \frac{\sum_{i=1}^n w_j(x_i)^f x_i}{\sum_{i=1}^n w_j(x_i)^f}$$



JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

© The Johns Hopkins University 2022, All Rights Reserved.





JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

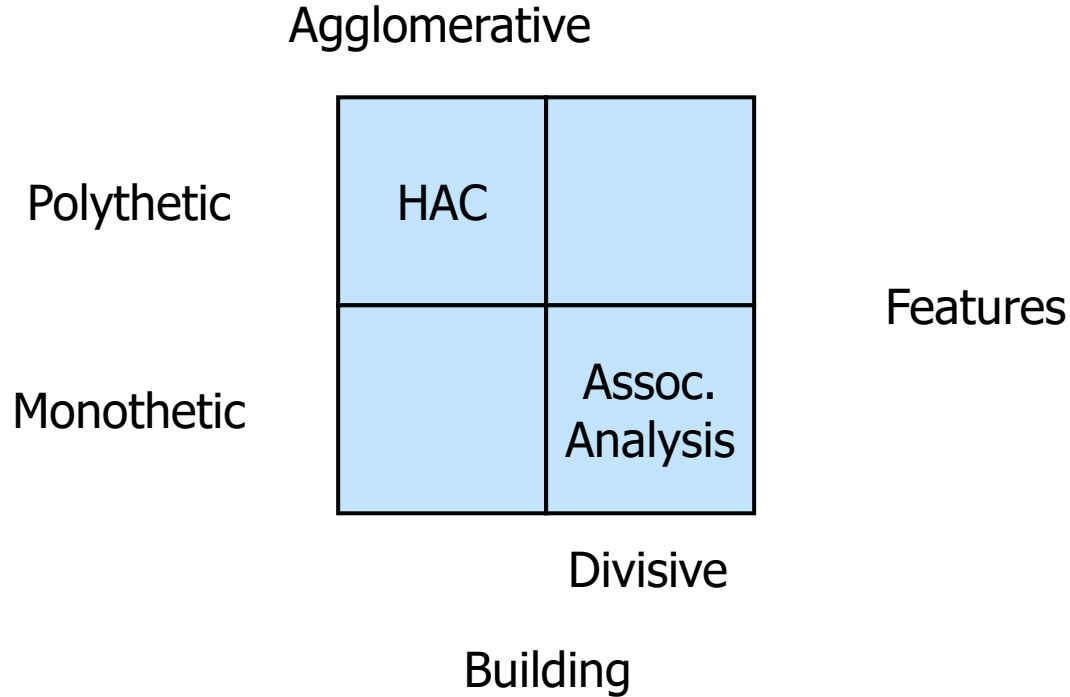
Introduction to Machine Learning

Hierarchical Agglomerative Clustering


Dendrogram



Agglomerative vs Divisive



Hierarchical Agglomerative Clustering



Association Analysis

- Recall

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}$$

- Binary \rightarrow Normalized
- F_{ij} is the j^{th} attribute of the i^{th} data point
- F_{ik} is the k^{th} attribute of the i^{th} data point

$$a_{jk} = \sum_{x_i} F_{ij} \times F_{ik}$$

$$b_{jk} = \sum_{x_i} (1 - F_{ij}) \times F_{ik}$$

$$c_{jk} = \sum_{x_i} F_{ij} \times (1 - F_{ik})$$

$$d_{jk} = \sum_{x_i} (1 - F_{ij}) \times (1 - F_{ik})$$

$$\chi_{jk}^2 = \frac{(ad - bc)^2}{(a - b)(a - c)(b - d)(c - d)}$$

$$F_{split} = \operatorname{argmax} \sum_{k=1}^d \chi_{jk}^2$$



JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

© The Johns Hopkins University 2022, All Rights Reserved.




JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

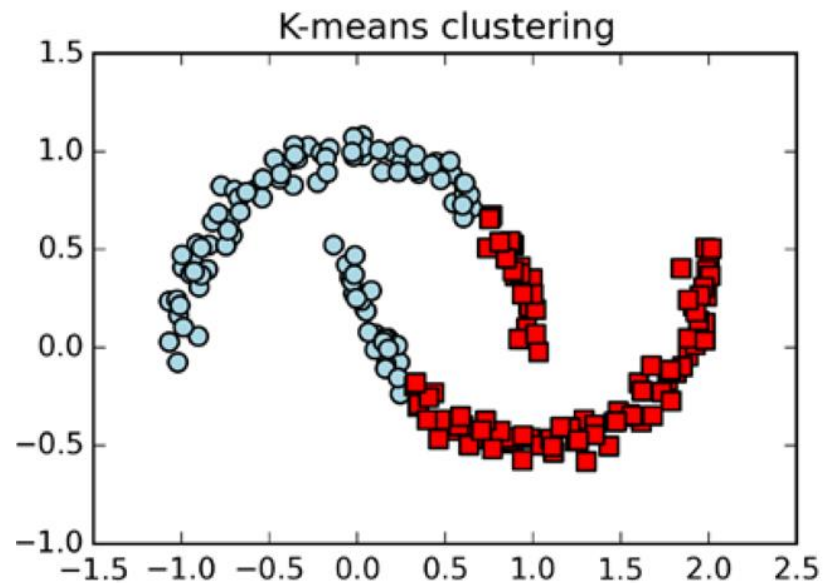
Introduction to Machine Learning


Density-based Clustering

A Motivating Example




K-Means






Density-based Clustering



Point Categorization



DBSCAN


Algorithm 2 DB-Scan

```
1: function DB-SCAN( $\mathcal{D}$ )
2:    $currClustLbl \leftarrow 1$ 
3:   for all  $p \in Core$  do do
4:     if  $clustLbl[p] = \text{"Unknown"}$  then
5:        $currClustLbl \leftarrow currClustLbl + 1$ 
6:        $clustLbl[p] \leftarrow currClustLbl$ 
7:     end if
8:     for all  $p' \in \theta\text{-neighborhood}$  do
9:       if  $clustLbl[p'] = \text{"Unknown"}$  then
10:         $clustLbl[p'] \leftarrow currClustLbl$ 
11:      end if
12:    end for
13:  end for
14:  return  $clustLbl$ 
15: end function
```

Another Example



Another Example





JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

© The Johns Hopkins University 2022, All Rights Reserved.



JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

Introduction to Machine Learning

Spectral Clustering

Laplacian Matrix

- Begin with a similarity matrix, $\mathbf{M} = \begin{bmatrix} \delta(x_1, x_1) & \cdots & \delta(x_1, x_n) \\ \vdots & \ddots & \vdots \\ \delta(x_n, x_1) & \cdots & \delta(x_n, x_n) \end{bmatrix}$
- The ϵ -neighborhood graph limits points for which $\delta(x_i, x_j)$ is calculated to require that $\delta(x_i, x_j) \leq \epsilon$.
- The k -nearest neighbor graph limits points corresponding to, as the name suggests, the k nearest neighbors of each point.
- These approaches lead to “reduced” similarity matrices, $\mathbf{M}_{\text{reduced}}$.
- The graph Laplacian begins with a diagonal “degree” matrix, $\mathbf{\Delta} = \begin{bmatrix} \text{deg}(x_1) & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \text{deg}(x_n) \end{bmatrix}$,
where $\text{deg}(x_i) = \sum_j \delta(x_i, x_j)$.
- Then the graph Laplacian is $\mathbf{L} = \mathbf{\Delta} - \mathbf{M}_{\text{reduced}}$.

Un-normalized Spectral Clustering


1. Construct $\mathbf{M}_{\text{reduced}}$.
2. Construct $\mathbf{L} = \mathbf{\Delta} - \mathbf{M}_{\text{reduced}}$.
3. Find the first k non-zero eigenvectors $\mathbf{u}_1, \dots, \mathbf{u}_k$ of \mathbf{L} .
4. Construct matrix \mathbf{U} from $\mathbf{u}_1, \dots, \mathbf{u}_k$.
5. Cluster the *rows* of \mathbf{U} with k -means.
6. Return the row indices grouped by the clusters.

Example Clusters


Data



Symmetric KNN



Spectral Clustering



Normalized (Sym) Spectral Clustering

Two normalized graph Laplacians

$$\mathbf{L}_{\text{sym}} = \Delta^{-1/2} \mathbf{L} \Delta^{-1/2} = \mathbf{I} - \Delta^{-1/2} \mathbf{U} \Delta^{-1/2}$$

$$\mathbf{L}_{\text{rw}} = \Delta^{-1/2} \mathbf{L} = \mathbf{I} - \Delta^{-1/2} \mathbf{U}$$

1. Construct $\mathbf{M}_{\text{reduced}}$.
2. Construct $\mathbf{L} = \Delta - \mathbf{M}_{\text{reduced}}$.
3. Construct $\mathbf{L}_{\text{sym}} = \Delta^{-1/2} \mathbf{L} \Delta^{-1/2}$.
4. Find the first k non-zero eigenvectors $\mathbf{u}_1, \dots, \mathbf{u}_k$ of \mathbf{L}_{sym} .
5. Construct matrix \mathbf{U} from $\mathbf{u}_1, \dots, \mathbf{u}_k$.
6. Normalize the rows of \mathbf{U} , $\forall i \leq n, \sum_j u_{ij}^2 = 1$.
7. Cluster the *rows* of \mathbf{U} with k -means.
8. Return the row indices grouped by the clusters.

Normalized (RW) Spectral Clustering

Two normalized graph Laplacians

$$\mathbf{L}_{\text{sym}} = \Delta^{-1/2} \mathbf{L} \Delta^{-1/2} = \mathbf{I} - \Delta^{-1/2} \mathbf{U} \Delta^{-1/2}$$

$$\mathbf{L}_{\text{rw}} = \Delta^{-1/2} \mathbf{L} = \mathbf{I} - \Delta^{-1/2} \mathbf{U}$$

1. Construct $\mathbf{M}_{\text{reduced}}$.
2. Construct $\mathbf{L} = \Delta - \mathbf{M}_{\text{reduced}}$.
3. Construct $\mathbf{L}_{\text{rw}} = \Delta^{-1/2} \mathbf{L}$.
4. Find the first k non-zero eigenvectors $\mathbf{u}_1, \dots, \mathbf{u}_k$ of \mathbf{L}_{sym} .
5. Construct matrix \mathbf{U} from $\mathbf{u}_1, \dots, \mathbf{u}_k$.
6. Cluster the *rows* of \mathbf{U} with k -means.
7. Return the row indices grouped by the clusters.



JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

© The Johns Hopkins University 2022, All Rights Reserved.



JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

Introduction to Machine Learning

PAC Learning

Sample Complexity


Computational Learning Theory (COLT)

1. What general laws constrain learning?
2. What types of learning problems can be solved in reasonable time/space?
3. When can we trust the output of a learned hypothesis, and by how much can we trust it?

Types of Supervised Learning

- Teacher-annotated learning
 - Probability distribution
 - Inductive learning principle/hypothesis
- Active learning
 - Examples \rightarrow oracle
 - Oracle provides the label
- Helpful teacher
 - Teacher picks the examples
 - Principle to determine minimum number of examples needed

Learning with a Helpful Teacher



Back to Sample Complexity

- The sample complexity of a learning problem seeks to determine, for a concept,
 - The number of training examples needed to learn the concept
 - Ideally minimize the amount of training required under the associated learning model
- The notion of sample complexity was first posed by Leslie Valiant.

L. G. Valiant, "A Theory of the Learnable," *Communications of the ACM*, Volume 27, Issue 11, November 1984, pp. 1134–1142.
- **Def:** A concept C is said to be *PAC-Learnable* by learner L using hypothesis space H if, for all $c \in C$, distributions D over X of length n , ε such that $0 < \varepsilon < 0.5$, and δ such that $0 < \delta < 0.5$, learner L will output hypothesis $h \in H$ with probability at least $(1 - \delta)$ such that the true error $error(h) \leq \varepsilon$ in time polynomial in $\frac{1}{\varepsilon}, \frac{1}{\delta}, n$ and $size(c)$.

PAC Learning

- Term first coined by Dana Angluin.
- Many key theorems proven by David Haussler.

$$m \geq \frac{1}{\epsilon} \left(\ln |H| + \ln \frac{1}{\delta} \right)$$

- Example

$$y = f(x_1, \dots, x_n) = x'_1 \wedge \dots \wedge x'_k$$

- Active learner: n examples
- Helpful teacher: k examples
- Teacher annotated: $\frac{1}{\epsilon} \left(\ln n + \ln \frac{1}{\delta} \right)$ examples



JOHNS HOPKINS

WHITING SCHOOL
of ENGINEERING

© The Johns Hopkins University 2022, All Rights Reserved.